

Week 7 Progress Report

Name: Blake Simpson Week: 7 (Polish + final report draft) Track: Track 3 — Agent

What I did this week

- Fixed the top 3 issues from the Week 6 eval:
 - i. **More jobs**: added two sources (We Work Remotely, Working Nomads): 241 -> ~710 postings.
 - ii. **No padding**: search returns nothing instead of filler, wrong-role jobs cap at a score of 3, and the CLI admits "nothing scored 7+ today."
 - iii. **Smarter stopping**: the agent quits hopeless searches instead of burning all 3 rounds.
- Re-ran the 12 eval cases: success 2/12 -> 8/12, precision@5 0.32 -> 0.65 (full tables in `docs/eval/results.md`).
- Cleanup: removed unused dependencies, fixed a job-id bug that could re-send alerts, rewrote the README, verified a fresh clone runs (25 tests pass).
- Drafted the final report (`docs/final-report.md`) and the Week 8 slide deck (`docs/week-8-deck.md`).

What worked

- Reusing the same 12 test cases made the before/after numbers directly comparable.
- The supply fix worked as predicted: winnable cases went from 2 to 8 and the agent passed all 8. The remaining fails just had too few relevant postings to pass — and the agent found every one that existed.
- The new stop rule fired exactly where intended: two hopeless cases stopped at round 2 instead of wasting a third.

What failed or surprised me

- My first stop rule (quit when two queries find nothing new) almost never fires on the bigger corpus — hopeless queries now find new-but-irrelevant jobs instead of nothing. The eval data pointed to a better rule (zero good

matches after two rounds -> stop) and killed a stricter one that would have cost python_backend a win.

- The no-Claude baseline also jumped (1/12 -> 7/12) just from the corpus fix — most of this week's gain was supply, as Week 6 predicted.

What I learned

- Fixing supply (adding RSS feeds) moved the numbers more than any agent-logic change all project - confirmation that evaluating before optimizing was the right call.

Evidence of progress

- New code: `src/job_scout/weworkremotely.py` , `src/job_scout/workingnomads.py` , `src/job_scout/ids.py` , `tests/test_week7_fixes.py` (25 tests total, passing)
- Changed: `jobs.py` (no padding), `agent.py` (exhaustion stop, score-only-new, stop_reason), `tools.py` (calibrated scoring prompt), `main.py` (honest output)
- Results: `docs/eval/results.md` / `results.json` (Week 7 re-run), `docs/evaluation.md` 9 (before/after)
- Report + deck: `docs/final-report.md` , `docs/week-8-deck.md`

Plan for next week

- Demo the final prototype

Blockers

- None